

S3-TASK-01

1. Create an S3 bucket and upload some objects to S3.


Step 1: Sign in to AWS Management Console

1. Open the AWS Management Console.
2. Search for **S3** in the search bar.
3. Click **Amazon S3**.

Step 2: Create an S3 Bucket

1. Click **Create bucket**.
2. Enter a **Bucket name** (must be globally unique).
 - Example: srihari-files
3. Select the **AWS Region**.
 - Example: Europe (Stockholm) eu-north-1
4. Leave the remaining settings as default (unless your lab specifies otherwise).
5. Click **Create bucket**.

Expected Result: The S3 bucket is created successfully.



Step 3: Open the Bucket

1. Click the newly created bucket name (srihari-files).
2. The **Objects** tab opens automatically.

Step 4: Upload Objects

1. Click **Upload**.
2. Click **Add files**.

3. Select one or more files from your computer.


○ Example:

- git-cheat-sheet.pdf
- image.jpg
- notes.txt

4. Click **Open**.

5. Click **Upload**.

Expected Result: The selected files are uploaded successfully.



Step 5: Verify the Uploaded Objects

1. Open the **Objects** tab of the bucket.
2. Verify that the uploaded files are listed.

Example:

Bucket Name: srihari-files

Objects

git-cheat-sheet.pdf

image.jpg

notes.txt


Step 6: Verify Using AWS CLI (Optional)

List all buckets:

```
aws s3 ls
```

List objects in your bucket:

```
aws s3 ls s3://srihari-files
```



NOTE: Make sure you enable the permissions (public ACL) , and uncheck the bucket permissions while creating

2. Deploy a static website in the S3 bucket.

Step 1: Create or Select an S3 Bucket

1. Sign in to the **AWS Management Console**.
2. Search for **Amazon S3** and open it.
3. Create a new bucket or select an existing bucket (e.g., srihari-files).

Step 2: Prepare Website Files

Save these files as **index.html** and **error.html**.

Step 3: Upload Website Files


1. Open your S3 bucket.
2. Go to the **Objects** tab.
3. Click **Upload**.
4. Click **Add files**.
5. Select:
 - index.html
 - error.html

6. Click **Upload**.

Step 4: Enable Static Website Hosting

1. Open the bucket.
2. Go to the **Properties** tab.
3. Scroll to **Static website hosting**.

4. Click **Edit**.
5. Select **Enable**.
6. Choose **Host a static website**.
7. Enter:
 - **Index document:** index.html
 - **Error document:** error.html
8. Click **Save changes**.



Step 5: Disable Block Public Access

1. Go to the **Permissions** tab.
2. Under **Block public access (bucket settings)**, click **Edit**.
3. Uncheck **Block all public access**.
4. Select the acknowledgment checkbox.

5. Click **Save changes**.

Step 6: Add a Bucket Policy

Go to **Permissions** → **Bucket Policy** and paste the following policy (replace srihari-files with your bucket name if different):

The screenshot shows the AWS IAM console interface for editing a bucket policy. The policy is a JSON document that grants 'PublicReadGetObject' permission to the 'arn:aws:s3::srihari-files/*' resource. The console also shows a summary of the upload operation and a table of files and folders.

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "PublicReadGetObject",
6       "Effect": "Allow",
7       "Principal": "*",
8       "Action": "s3:GetObject",
9       "Resource": "arn:aws:s3::srihari-files/*"
10    }
11  ]
12 }
```

Summary

Destination	Succeeded	Failed
s3://srihari-files	2 files, 35.0 B (100.00%)	0 files, 0 B (0%)

Files and folders (2 total, 35.0 B)

Name	Folder	Type	Size	Status	Error
error.html	-	text/html	17.0 B	Succeeded	-
maresh.html	-	text/html	18.0 B	Succeeded	-

Step 7: Access the Website

1. Go to the **Properties** tab.
2. Scroll to **Static website hosting**.
3. Copy the **Bucket website endpoint**.
4. Open the endpoint in a web browser.

Example:

<http://srihari-files.s3-website-eu-north-1.amazonaws.com>


Step 8: Verify the Website

- The browser should display:

Welcome to My S3 Static Website

This website is hosted using Amazon S3.


If you open a page that doesn't exist, S3 will display the error.html page.



3. Enable cross-region replication on S3 buckets.

Step 1: Create the Source Bucket


1. Sign in to the AWS Management Console.
2. Search for **S3**.
3. Click **Create bucket**.
4. Enter a bucket name.
 - Example: my-source-bucket123
5. Select Region:
 - **Asia Pacific (Mumbai) – ap-south-1**
6. Keep the default settings.
7. Click **Create bucket**.



Step 2: Create the Destination Bucket

1. Click **Create bucket** again.
2. Enter a different bucket name.
 - Example: my-destination-bucket123

3. Select another Region.
 - **Asia Pacific (Singapore) – ap-southeast-1**
4. Click **Create bucket**.



Step 3: Enable Versioning on Both Buckets

Source Bucket

1. Open the source bucket.
2. Go to the **Properties** tab.
3. Scroll to **Bucket Versioning**.
4. Click **Edit**.
5. Select **Enable**.
6. Click **Save changes**.

aws [Search] [Alt+S] Europe (Stockholm) Srihari (299582146630) Srihari

Amazon S3 > Buckets > my-source-bucket123-299582146630-eu-north-1-an

Amazon S3

- ▼ Buckets
 - General purpose buckets
 - Directory buckets
 - Table buckets
 - Vector buckets
- ▼ Files
 - File systems
- ▼ Access management and security
 - Access Points
 - Access Points for FSx
 - Access Grants
 - IAM Access Analyzer

Successfully edited Bucket Versioning
To transition, archive, or delete older object versions, [configure lifecycle rules](#) for this bucket.

my-source-bucket123-299582146630-eu-north-1-an Info

< Objects Metadata **Properties** Permissions Metrics >

Bucket overview

AWS Region Europe (Stockholm) eu-north-1	Amazon Resource Name (ARN) arn:aws:s3:::my-source-bucket123-299582146630-eu-north-1-an
Creation date July 21, 2026, 19:36:35 (UTC+05:30)	

Destination Bucket

Repeat the same steps:

- Open the destination bucket.
- Go to **Properties**.
- Enable **Bucket Versioning**.
- Save the changes.

Note: CRR will not work unless versioning is enabled on both buckets.

aws [Search] [Alt+S] United States (N. Vir) Srihari (299582146630) Srihari

Amazon S3 > Buckets > my-destination-bucket123-299582146630-us-east-1-an

Successfully edited Bucket Versioning
To transition, archive, or delete older object versions, [configure lifecycle rules](#) for this bucket.

my-destination-bucket123-299582146630-us-east-1-an Info

< Objects Metadata **Properties** Permissions Metrics Management File systems - n >

Bucket overview

AWS Region US East (N. Virginia) us-east-1	Amazon Resource Name (ARN) arn:aws:s3:::my-destination-bucket123-299582146630-us-east-1-an	Creation date July 21, 2026, 20:01:30 (UTC+05:30)
--	--	---

Step 4: Configure Replication

1. Open the **Source Bucket**.
2. Go to the **Management** tab.

3. Scroll to **Replication rules**.
4. Click **Create replication rule**.

Step 5: Configure the Rule

Rule Name

Example:

CRR-Rule

Rule Status

Select:

Enabled

Choose Objects

You can replicate:

- All objects (recommended for practice)
- Specific folder (prefix)
- Objects with specific tags

Select:

Apply to all objects

Click **Next**.

aws

Search

[Alt+S]

Europe (Stockholm)

Srihari (299582146630)

Srihari

Amazon S3

Replication rules

Create replication rule

Create replication rule [Info](#)

Use replication rules to define options you want Amazon S3 to apply during replication such as server-side encryption, replica ownership, transitioning replicas to another storage class, and more. Replication also includes any annotations attached to your objects. [Learn more](#)

Replication now includes replicating any annotations attached to your objects, which might increase your costs. To understand the potential cost impact, review your replication rules and annotations usage.

Replication rule configuration

Replication rule name

CRR-Rule

Up to 255 characters. In order to be able to use CloudWatch metrics to monitor the progress of your replication rule, the replication rule name must only contain English characters.

Status

Choose whether the rule will be enabled or disabled when created.

☒ Enabled

☐ Disabled

Priority

The priority value resolves conflicts that occur when an object is eligible for replication under multiple rules to the same destination. The rule is added to the configuration at the highest priority and the priority can be changed on the replication rules table.

0

aws

Search

[Alt+S]

Europe (Stockholm)

Srihari (299582146630)

Srihari

Amazon S3

Buckets

my-source-bucket123-299582146630-eu-north-1-an

Replication rules

Replication configuration successfully updated

If changes to the configuration aren't displayed, choose the refresh button. Changes apply only to new objects. To replicate existing objects with this configuration, choose **Create replication job**.

Create replication job

Replication rules [Info](#)

Replication enables automatic and asynchronous copying of objects across buckets in the same or different AWS Regions. A replication configuration is a set of rules that define what options should be applied to a group of objects during replication.

Replication now includes replicating any annotations attached to your objects, which might increase your costs. To understand the potential cost impact, review your replication rules and annotations usage.

Replication configuration settings

Configuration settings affect all replication rules in the bucket.

Source bucket

my-source-bucket123-299582146630-eu-north-1-an

Source Region

Europe (Stockholm) eu-north-1

IAM role

[s3crr_role_for_my-source-bucket123-299582146630-eu-north-1-an](#)

Create replication job

Edit

Step 6: Select the Destination Bucket

Choose:

Destination bucket

Select:

Choose a bucket in this account

Choose your destination bucket.

Example:

my-destination-bucket123

Step 7: IAM Role

AWS asks for an IAM Role.

Choose:

Create new IAM role

AWS creates the required role automatically.

Step 8: Replication Options

Keep the default settings unless your task requires changes.

Examples:

- Replicate delete markers
- Replicate existing objects (optional)
- Storage class

Click **Save**.

Your replication rule is now created.

Step 9: Test Cross-Region Replication

1. Open the **Source Bucket**.
2. Click **Upload**.
3. Upload a file.

Example:

hello.txt

Wait a few seconds.

Upload [Info](#)

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDKs or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files** or **Add folder**.

Files and folders (1 total, 169.2 KB)

[Remove](#)[Add files](#)[Add folder](#)

All files and folders in this table will be uploaded.

< 1 >

☒	Name	Folder	Type	Size
☒	BOORA SRIHARI RESUME....	-	application/pdf	169.2 KB

Destination [Info](#)

Destination
<s3://my-source-bucket123-299582146630-eu-north-1-an>

Upload succeeded

For more information, see the **Files and folders** table.

Upload: status [Close](#)

After you navigate away from this page, the following information is no longer available.

Summary

Destination
<s3://my-source-bucket123-299582146630-eu-north-1-an>

Succeeded
 1 file, 169.2 KB (100.00%)

Failed
 0 files, 0 B (0%)

[Files and folders](#) [Configuration](#)

Files and folders (1 total, 169.2 KB)

< 1 >


Name	Folder	Type	Size	Status	Error
BOORA SRIHARI R...	-	application/pdf	169.2 KB	Succeeded	-

Step 10: Verify Replication

Open the **Destination Bucket**.

You should see:

hello.txtThe object has been automatically replicated.



4. Configure a bucket policy so only the Admin user can see the objects of the S3 bucket.

Step 1: Find the Admin User ARN

1. Open **AWS Console**.
2. Search for **IAM**.
3. Click **Users**.
4. Click on the **Admin** user.
5. Copy the **User ARN**.

Example:

arn:aws:iam::123456789012:user/Admin

Replace this with your own Admin user's ARN.

Step 2: Open the S3 Bucket

1. Open **S3**.
2. Select your bucket.
3. Go to the **Permissions** tab.
4. Scroll to **Bucket Policy**.
5. Click **Edit**.


Step 3: Add the Bucket Policy

Replace:

- 123456789012 with your **AWS Account ID**

- your-bucket-name with your **bucket name**


Click **Save changes**.



Step 4: Test

- Log in as the **Admin** user.
- Open the bucket.
- You should be able to:
 - ☒ View the bucket
 - ☒ List the objects
 - ☒ Download objects

A different IAM user without permission should **not** be able to access the bucket.



5. Set up lifecycle policies to automatically transition or delete objects based on specific criteria.

Step 1: Open Amazon S3

1. Log in to the **AWS Management Console**.
2. Search for **S3**.
3. Click **Amazon S3**.

Step 2: Open Your Bucket

1. Click the bucket where you want to create the lifecycle policy.
 - Example: srihari-bucket

Step 3: Go to the Management Tab


At the top of the bucket page, click:

Objects | Properties | Permissions | Metrics | Management

Click **Management**.

Step 4: Create a Lifecycle Rule

1. Under **Lifecycle rules**, click **Create lifecycle rule**.



Step 5: Enter Rule Details

Rule Name

Example:

Delete-Old-Files

or

Lifecycle-Rule

Choose Rule Scope

Select:

☒ **Apply to all objects in the bucket**

AWS will display a warning.

Check the confirmation box.

Click **Next**.

Note: If you want the rule to apply only to specific objects, you can use a **prefix** (folder name) or **object tags** instead.

aws [Search] [Alt+S] Europe (Stockholm) Srihari (299582146630) srihari

Amazon S3 > Buckets > srihari-files > Lifecycle configuration > Create lifecycle rule

Create lifecycle rule [Info](#)

Lifecycle rule configuration

Lifecycle rule name

Delete-Old-Files

Up to 255 characters

Choose a rule scope

☐ Limit the scope of this rule using one or more filters

☒ Apply to all objects in the bucket

⚠ Apply to all objects in the bucket

If you want the rule to apply to specific objects, you must use a filter to identify those objects. Choose "Limit the scope of this rule using one or more filters". [Learn more](#)

☒ I acknowledge that this rule will apply to all objects in the bucket.

Step 6: Configure Lifecycle Actions

You can choose one or more actions.

Option 1: Transition Objects (Move to a Lower-Cost Storage Class)

Check:

☒ **Transition current versions of objects between storage classes**

Enter:

After: 30 Days

Move to: Standard-IA

Or

After: 90 Days

Move to: Glacier Flexible Retrieval

Option 2: Expire (Delete) Objects

Check:

☒ **Expire current versions of objects**

Enter:

Delete after: 365 Days

AWS will automatically delete the objects after 365 days.

Example Configuration

Action	Time
Transition to Standard-IA	30 Days
Transition to Glacier	90 Days
Delete Object	365 Days

Review transition and expiration actions

Current version actions

Day 0

- Objects uploaded

↓

Day 30

- Objects move to Standard-IA

↓

Day 365

- Objects expire

Noncurrent versions actions

Day 0

No actions defined.

Cancel>Create rule

aws

Search

[Alt+S]

Europe (Stockholm)

Srihari (299582146630)

srihar

Amazon S3

Buckets

srihari-files

Lifecycle configuration

✔ The rule "Delete-Old-Files" has been successfully added and the lifecycle configuration has been updated

It may take some time for the configuration to be updated. Refresh the lifecycle rules list if changes to the configuration aren't displayed.

×

Lifecycle configuration

To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Default minimum object size for transitions

All storage classes 128K

Lifecycle rules (1)

↻

View details

Edit

Delete

Actions ▾

Create lifecycle rule

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

Find lifecycle rules by name

< 1 >

⚙

Lifecycl...	Status	Scope	Current ...	Noncurr...	Expired ...	Incomplet...
☐ Delete-Old-Files	✔ Enabled	Entire bucket	Transition to Sta	-	-	-

Step 7: Review the Rule

Review the settings.

Example:

Rule Name:

Lifecycle-Rule

Scope:

All Objects

Transition:

30 Days → Standard-IA

Expiration:

365 Days → Delete

Step 8: Create the Rule

Click:

Create Rule

The lifecycle policy is now active.


Step 9: Verify the Rule

1. Go to the **Management** tab.
2. Under **Lifecycle Rules**, you should see:

Lifecycle-Rule

Status: Enabled

This confirms the rule has been created successfully.



The screenshot shows the AWS Management Console interface for the 'srihari-files' bucket. The 'Management' tab is active, displaying the 'Lifecycle configuration' section. The page title is 'srihari-files'. The navigation bar includes tabs for 'adata', 'Properties', 'Permissions', 'Metrics', 'Management' (selected), 'File systems - new', and 'Access Points'. The 'Lifecycle configuration' section explains that lifecycle rules define actions for objects based on their age. It notes the 'Default minimum object size for transitions' is 128K. Below this, the 'Lifecycle rules (1)' section shows a table with one rule: 'Delete-Old-Files'. The rule is 'Enabled' and applies to the 'Entire bucket'. The action is 'Transition to Standard-IA' after 30 days. A 'View lifecycle configuration' link is at the bottom.

Lifecycle ru...	Status	Scope	Current ver...	Noncurrent...	Expired obj...	Incomplet...
Delete-Old-Files	Enabled	Entire bucket	Transition to Sta	-	-	-

How It Works

Suppose you upload a file today.

resume.pdf

AWS automatically performs the following actions:

Day 1

resume.pdf

|

Day 30

Moved to Standard-IA

|

Day 90

Moved to Glacier

|

Day 365

Deleted

No manual action is required.

6. Push some objects to S3 using the AWS CLI.

Step 1: Verify AWS CLI Installation

Open **Command Prompt** (Windows) or **Terminal** (Linux).

Run:

aws --version

Example Output

aws-cli/2.31.0 Python/3.13.7 Windows/10 exe/AMD64

If you see the version, AWS CLI is installed successfully.

Step 2: Configure AWS CLI

Run:

aws configure

Enter the following details:

AWS Access Key ID: XXXXXXXXXXXXXXXXXXXX

AWS Secret Access Key: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

Default region name: ap-south-1

Default output format: json

Note: You can get the Access Key ID and Secret Access Key from **IAM → Users → Security credentials**.

Step 3: Check Whether Configuration is Successful

Run:

```
aws s3 ls
```

Example Output

```
2026-07-21 10:45:20 srihari-bucket
```

```
2026-07-21 11:15:42 backup-bucket
```

This command lists all the S3 buckets in your AWS account.

Step 4: Create a Sample File

Create a text file on your computer.

Example:

```
test.txt
```

Content:

```
Hello AWS S3 CLI
```

Step 5: Upload the File to S3

Syntax

```
aws s3 cp <local-file> s3://<bucket-name>/
```

Example

If your bucket name is:

```
srihari-bucket
```

Run:

```
aws s3 cp test.txt s3://srihari-bucket/
```

Step 6: Verify Upload

After running the command, you should see:

```
upload: .\test.txt to s3://srihari-bucket/test.txt
```

This means the file has been uploaded successfully.

```

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws --version
aws-cli/2.35.23 Python/3.14.6 Windows/11 exe/AMD64

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws configure
AWS Access Key ID [*****PLHX]: AKIAULQD4ZBDORDKPLHX
AWS Secret Access Key [*****UXRS]: hBhnzwcBmmk9yfeDfpdmP6aog5o3hjFfsf
loUXRS
Default region name [eu-north-1]: eu-north-1
Default output format [json]: json

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws s3 ls
2026-07-21 20:01:30 my-destination-bucket123-299582146630-us-east-1-an
2026-07-21 19:36:35 my-source-bucket123-299582146630-eu-north-1-an
2026-07-21 20:56:58 srihari-files
2026-07-21 16:16:02 vpc-flowlogs-srihari

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ vi test.txt

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws s3 cp test.txt s3://srihari-files/
upload: .\test.txt to s3://srihari-files/test.txt

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ |

```

Step 7: Verify in AWS Console

1. Open **AWS Console**.
2. Go to **Amazon S3**.
3. Open your bucket.
4. You should see:

test.txt

The screenshot shows the AWS S3 console interface for the bucket 'srihari-files'. The 'Objects (1/4)' section displays a list of four objects. The 'test.txt' object is selected, highlighted in blue. The console includes navigation tabs for 'Amazon S3', 'Buckets', and 'srihari-files'. Action buttons like 'Copy S3 URI', 'Copy URL', 'Download', 'Open', 'Delete', and 'Actions' are visible. A search bar and 'Show versions' toggle are also present.

	Name	Type	Last modified	Size	Storage class
☐	error.html	html	July 21, 2026, 19:27:53 (UTC+05:30)	17.0 B	Standard
☐	git-cheat-sheet.pdf	pdf	July 21, 2026, 18:14:42 (UTC+05:30)	77.9 KB	Standard
☐	mahesh.html	html	July 21, 2026, 19:27:54 (UTC+05:30)	18.0 B	Standard
☒	test.txt	txt	July 21, 2026, 21:08:37 (UTC+05:30)	18.0 B	Standard

7. Write a Bash script to create an S3 bucket.

Check the AWS CLI:

```
aws --version
```

Bash Script

Create a file named **create_bucket.bash**.

```
#!/bin/bash

BUCKET_NAME="srihari-bucket-2026"
REGION="ap-south-1"

aws s3api create-bucket \
  --bucket $BUCKET_NAME \
  --region $REGION \
  --create-bucket-configuration LocationConstraint=$REGION

if [ $? -eq 0 ]; then
  echo "S3 Bucket '$BUCKET_NAME' created successfully."
else
  echo "Bucket '$BUCKET_NAME' already exists or an error occurred."
fi

~
~
~
```

Give Execute Permission

```
chmod +x create_bucket.sh
```

Run the Script

```
./create_bucket.sh
```

Expected Output

S3 Bucket 'srihari-bucket-2026' created successfully.

Verify the Bucket

List all S3 buckets:

```
aws s3 ls
```

Example Output:

2026-07-21 10:45:20 srihari-bucket-2026

```

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ pwd
/c/Users/SRIHARI/Downloads

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws configure
AWS Access Key ID [*****PLHX]: AKIAULQD4ZBDORDKPLHX
AWS Secret Access Key [*****UXRS]: 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
Default region name [eu-north-1]: eu-north-1
Default output format [json]: json

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ vi create_bucket01.bash

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ chmod +x create_bucket01.bash

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ ./create_bucket01.bash
{
  "Location": "http://srihari-bucket-2026.s3.amazonaws.com/",
  "BucketArn": "arn:aws:s3:::srihari-bucket-2026"
}
S3 Bucket 'srihari-bucket-2026' created successfully.

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws s3 ls
2026-07-21 20:01:30 my-destination-bucket123-299582146630-us-east-1-an
2026-07-21 19:36:35 my-source-bucket123-299582146630-eu-north-1-an
2026-07-21 21:23:56 srihari-bucket-2026
2026-07-21 20:56:58 srihari-files
2026-07-21 16:16:02 vpc-flowlogs-srihari

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ |

```

[Alt+S]
Europe (Stockholm) ▼
Srihari (299582146630) ▼
srihari

Amazon S3 > Buckets

Buckets

General purpose buckets
All AWS Regions
Directory buckets

General purpose buckets (1/5)
Info
Refresh
Copy ARN
Empty
Delete
Create bucket

Buckets are containers for data stored in S3.

< 1 >

	Name ▲	AWS Region ▼	Creation date ▼
☐	my-destination-bucket123-299582146630-us-east-1-an	US East (N. Virginia) us-east-1	July 21, 2026, 20:01:30 (UTC+05:30)
☐	my-source-bucket123-299582146630-eu-north-1-an	Europe (Stockholm) eu-north-1	July 21, 2026, 19:36:35 (UTC+05:30)
☒	srihari-bucket-2026	Asia Pacific (Mumbai) ap-south-1	July 21, 2026, 21:23:56 (UTC+05:30)
☐	srihari-files	Europe (Stockholm) eu-north-1	July 21, 2026, 18:09:29 (UTC+05:30)
☐	vpc-flowlogs-srihari	Europe (Stockholm) eu-north-1	July 16, 2026, 16:12:51 (UTC+05:30)

8. Upload a 1 GB file to S3 using the CLI.

Step 1: Verify AWS CLI

Open **Git Bash** or **Command Prompt** and run:

```
aws --version
```

Example Output:

```
aws-cli/2.31.0 Python/3.13.7 Windows/10 exe/AMD64
```

Step 2: Create a 1 GB File

Option 1: Linux / Git Bash

```
dd if=/dev/zero of=1GBfile.bin bs=1M count=1024
```

This creates a 1 GB file filled with zeros.

Step 3: Verify the File

Go to the folder or run:

```
ls -lh
```

Example:

```
-rw-r--r-- 1 user user 1.0G 1GBfile.bin
```

Step 4: Upload the File to S3

Syntax

```
aws s3 cp <file-name> s3://<bucket-name>/
```

Example

If your bucket name is:

```
srihari-files
```

Run:

```
aws s3 cp 1GBfile.bin s3://srihari-files/
```

Step 5: Wait for Upload

The CLI displays progress similar to:

```
Completed 256.0 MiB/1.0 GiB
```

```
Completed 512.0 MiB/1.0 GiB
```

```
Completed 768.0 MiB/1.0 GiB
```

```
Completed 1.0 GiB/1.0 GiB
```

Step 6: Verify Upload

Run:

aws s3 ls s3://srihari-files

Example Output:

2026-07-21 15:30:10 1073741824 1GBfile.bin

You can also:

1. Open **AWS Console**.
2. Go to **Amazon S3**.
3. Open your bucket.
4. Verify that 1GBfile.bin is present.

```
SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws --version
aws-cli/2.35.23 Python/3.14.6 Windows/11 exe/AMD64

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws configure
AWS Access Key ID [*****PLHX]: AKIAULQD4ZBDORDKPLHX
AWS Secret Access Key [*****UXRS]: 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
Default region name [eu-north-1]: eu-north-1
Default output format [json]: json

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ dd if=/dev/zero of=1GBfile.bin bs=1M count=1024
1024+0 records in
1024+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 1.58685 s, 677 MB/s

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ ls -lh
total 14G
-rw-r--r-- 1 SRIHARI 197121 159M Jun 23 14:45 02-Client-Server-Architecture.mp4
-rw-r--r-- 1 SRIHARI 197121 249M Jun 24 12:17 03-Linux-basics.mp4
-rw-r--r-- 1 SRIHARI 197121 225M Jun 25 11:39 04-Linux-basics.mp4
-rw-r--r-- 1 SRIHARI 197121 365M Jun 26 11:39 05-Linux-basics.mp4
-rw-r--r-- 1 SRIHARI 197121 396M Jun 30 11:58 '06-AppServers-and-Web Servers.mp4'
-rw-r--r-- 1 SRIHARI 197121 157M Jul 2 11:38 07-Bash-Scripting_Basics.mp4
-rw-r--r-- 1 SRIHARI 197121 96M Jul 3 12:06 08-Bash-Scripting_Basics.mp4
-rw-r--r-- 1 SRIHARI 197121 354M Jul 7 12:17 09-Git-and-Github.mp4
-rw-r--r-- 1 SRIHARI 197121 346M Jul 8 11:02 10-Git-and-Github.mp4
-rw-r--r-- 1 SRIHARI 197121 86M Jul 6 2024 '101 ESSAY Disha Publication (1).pdf'
-rw-r--r-- 1 SRIHARI 197121 86M Jul 6 2024 '101 ESSAY Disha Publication.pdf'
-rw-r--r-- 1 SRIHARI 197121 844K Aug 10 2024 '10nth naresh.jpeg'
-rw-r--r-- 1 SRIHARI 197121 356K Aug 10 2024 '10nth naresh_11zon.jpeg'
-rw-r--r-- 1 SRIHARI 197121 215M Jul 13 12:08 11-Introduction_to_Cloud.mp4
-rw-r--r-- 1 SRIHARI 197121 490M Jul 14 11:35 12-AWS-EC2_and_Azure-VM.mp4
-rw-r--r-- 1 SRIHARI 197121 406M Jul 16 13:17 '14-NAT_VPC-Flow-Logs_and_VPC-Peering (1).mp4'
-rw-r--r-- 1 SRIHARI 197121 237M Jul 17 12:07 15-VPC-Endpoint_and_Transit-Gateway.mp4
-rw-r--r-- 1 SRIHARI 197121 12M Jul 6 2024 '151 Essays for IAS _Main.pdf'
-rw-r--r-- 1 SRIHARI 197121 226M Jul 20 11:49 16-IAM_and_EC2-Snapshot_and_EBS-Volumes.mp4
-rw-r--r-- 1 SRIHARI 197121 205M Jul 21 11:49 17-S3-Bucket.mp4
-rw-r--r-- 1 SRIHARI 197121 0 Jul 22 11:36 '18-S3_and_Route53_and_CDN (1).mp4'
-rw-r--r-- 1 SRIHARI 197121 307M Jul 22 11:34 18-S3_and_Route53_and_CDN.mp4
-rw-r--r-- 1 SRIHARI 197121 1.0G Jul 22 11:44 1GBfile.bin

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws s3 cp 1GBfile.bin s3://srihari-files/
upload: .\1GBfile.bin to s3://srihari-files/1GBfile.bin

SRIHARI@LAPTOP-N3TTSS9I MINGW64 ~/Downloads
$ aws s3 ls s3://srihari-files
2026-07-22 11:49:04 1073741824 1GBfile.bin
2026-07-21 19:27:53 17 error.html
2026-07-21 18:14:42 79730 git-cheat-sheet.pdf
2026-07-21 19:27:54 18 mahesh.html
2026-07-21 21:08:37 18 test.txt
```